

Transformers vs vLLM: does the backend change the measured outcome?
60 paired prompts per model; same revision, rendered prompt, BF16, greedy decoding

